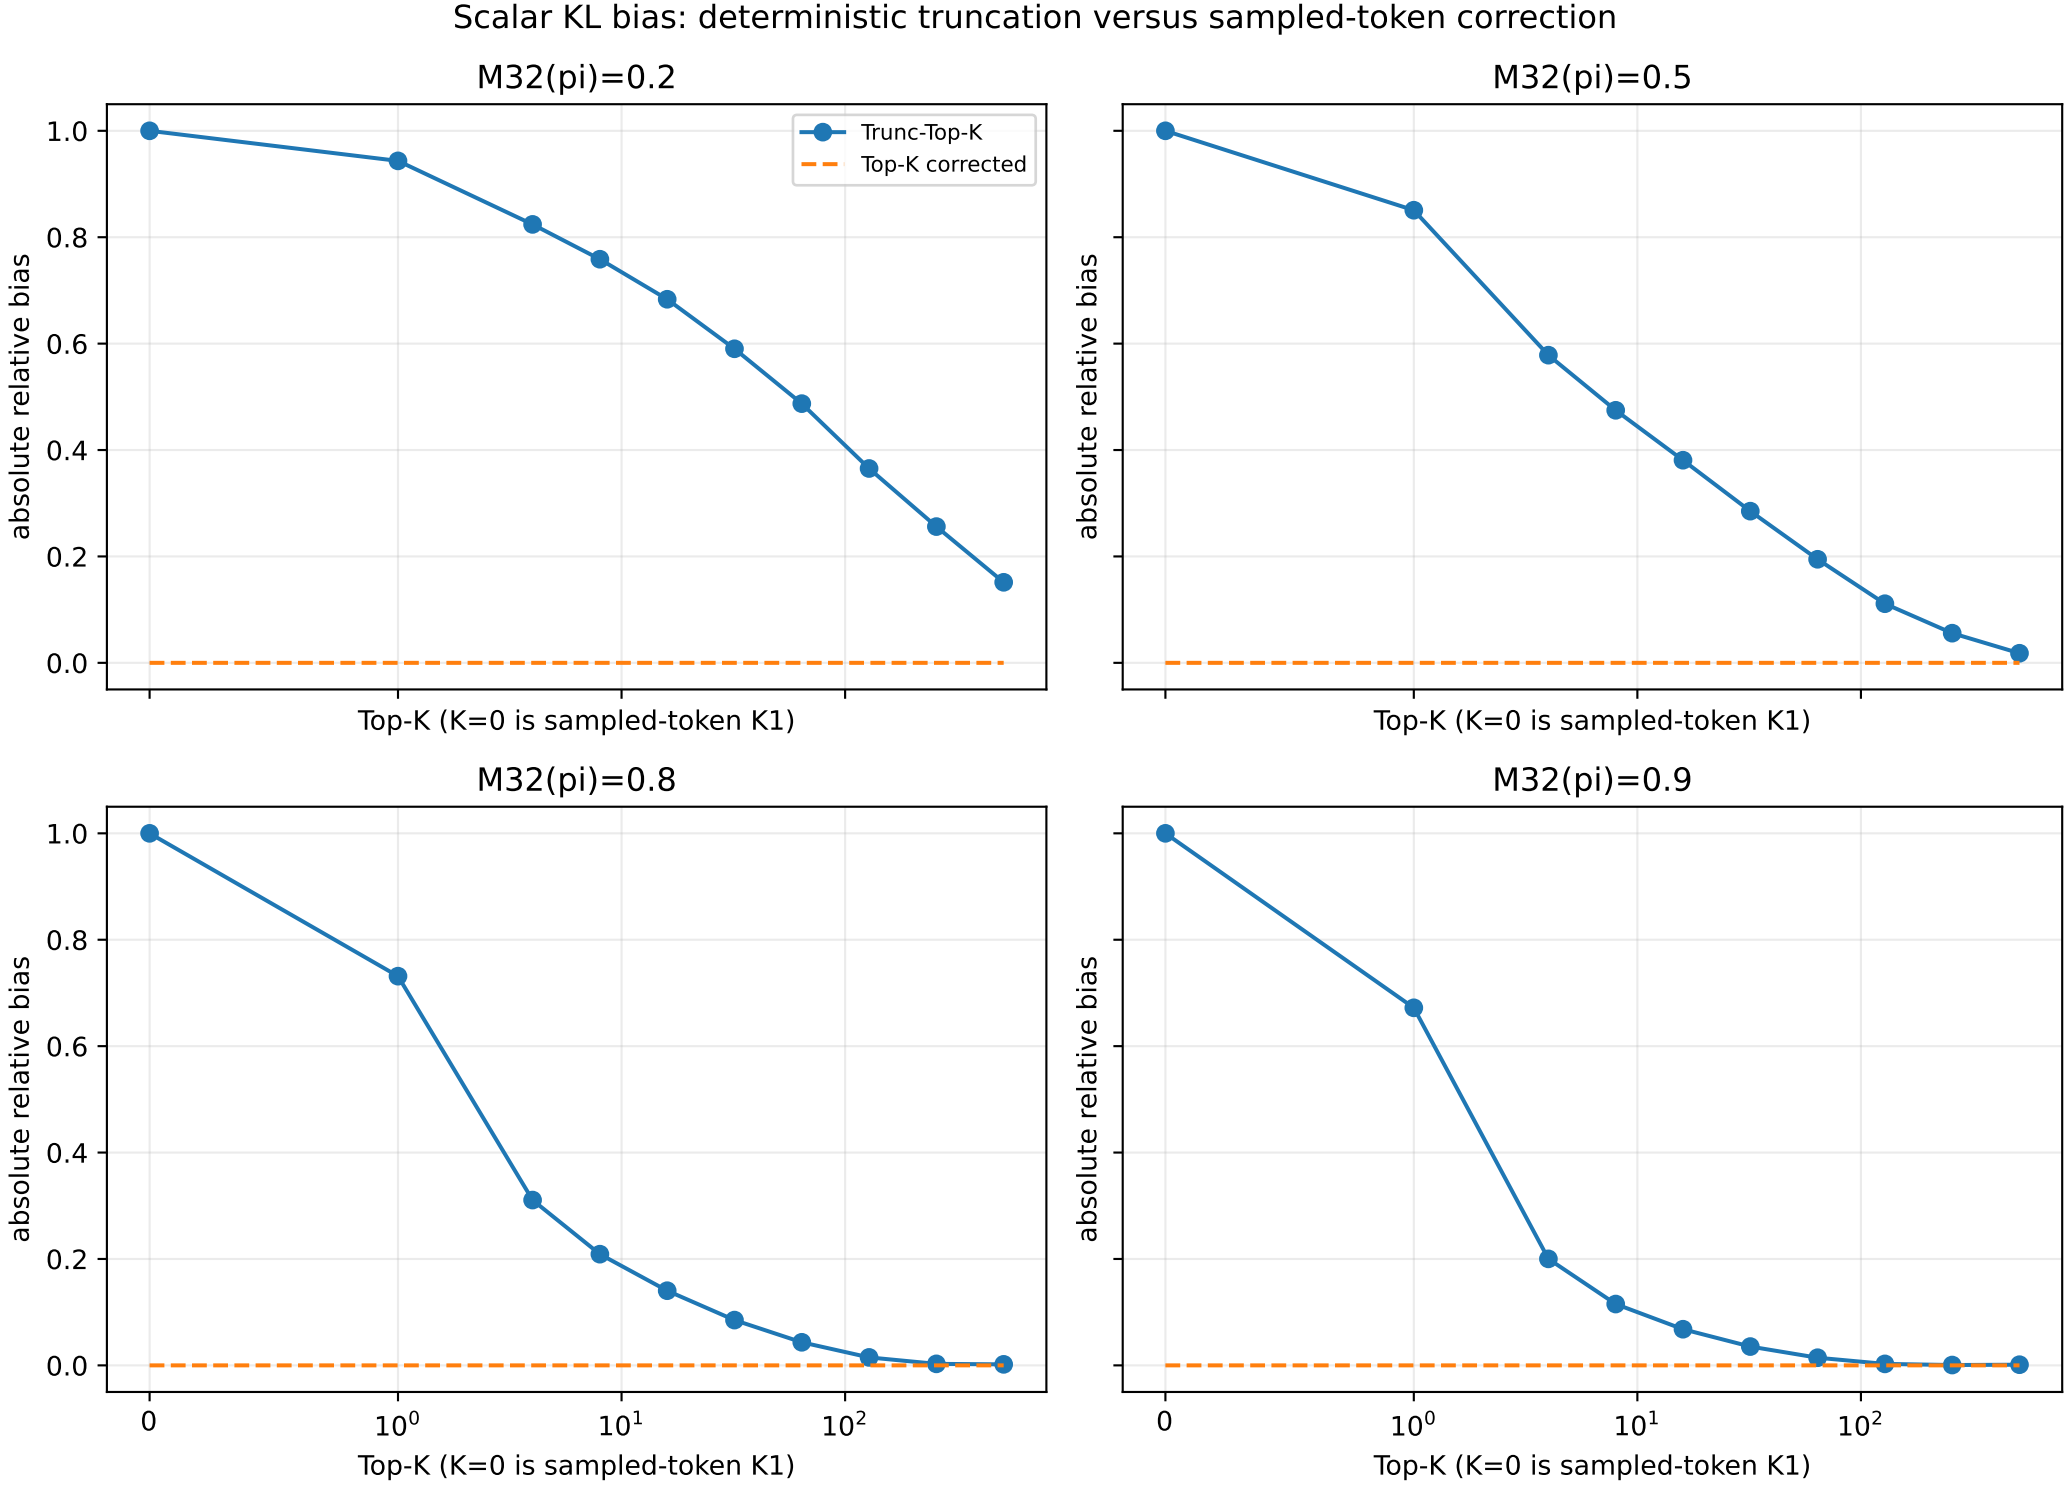
Scalar KL bias: deterministic truncation versus sampled-token correction
M32(pi)=0.2
M32(pi)=0.5
M32(pi)=0.8
M32(pi)=0.9
Trunc-Top-K
Top-K corrected
absolute relative bias
Top-K (K=0 is sampled-token K1)
0.0
0.2
0.4
0.6
0.8
1.0
0
10^0
10^1
10^2