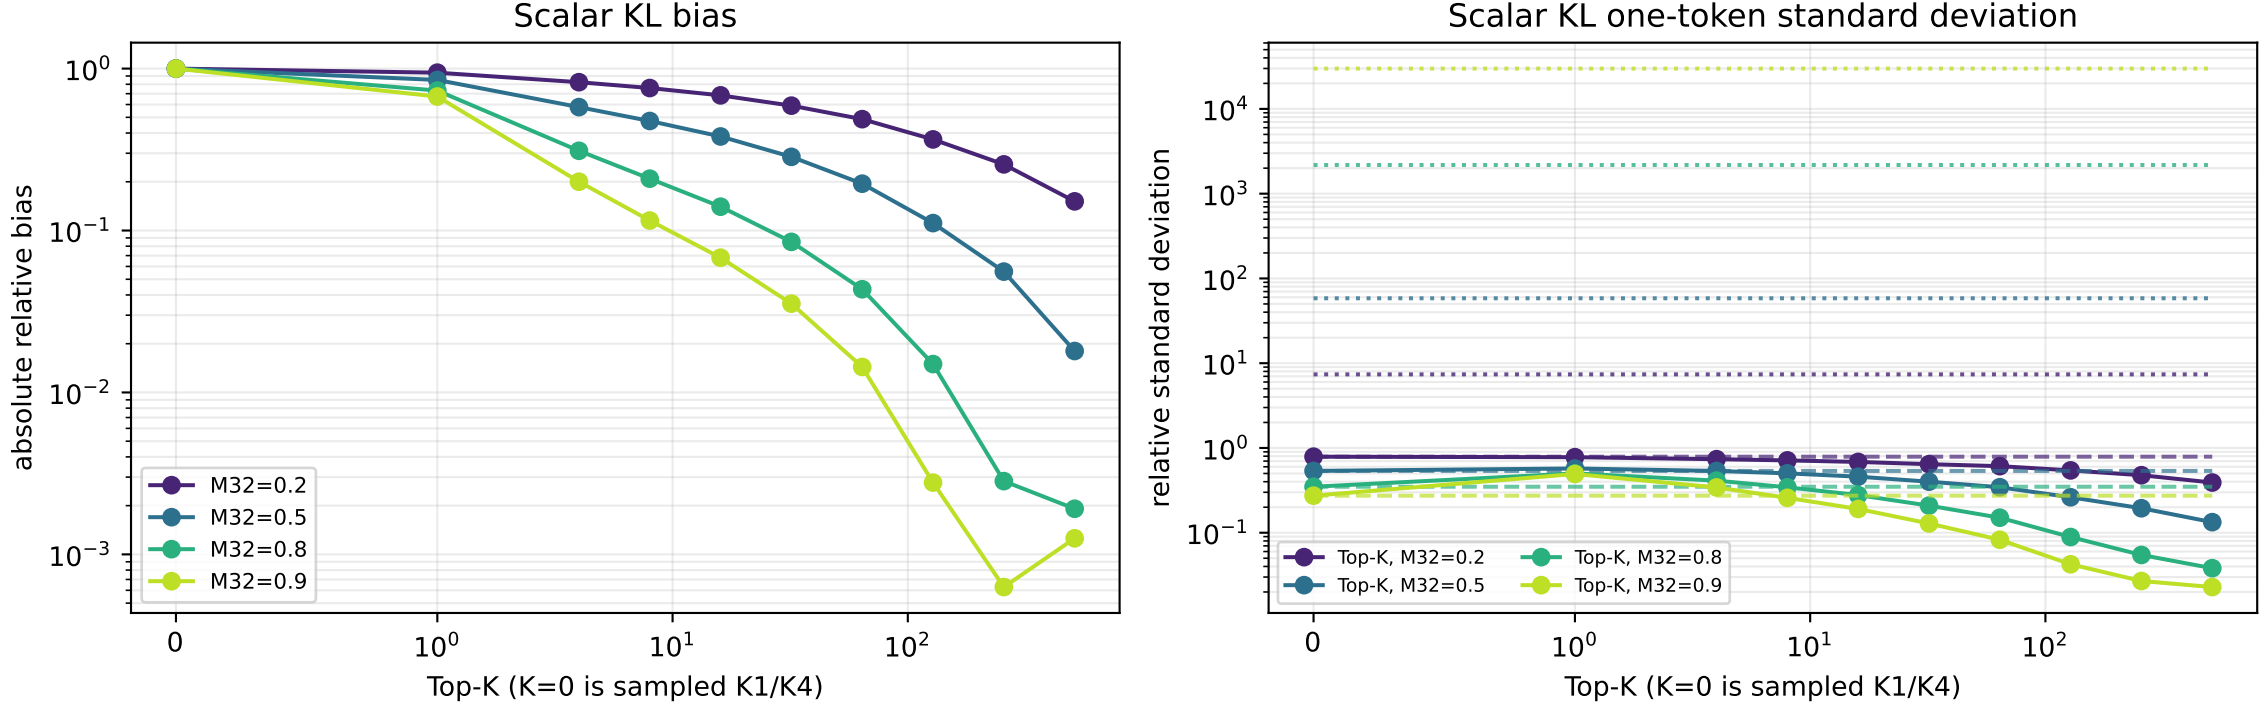
Scalar reverse-KL estimator: bias and variance
solid Top-K; dashed K1; dotted K3++
Scalar KL bias
absolute relative bias
Top-K (K=0 is sampled K1/K4)
M32=0.2
M32=0.5
M32=0.8
M32=0.9
Scalar KL one-token standard deviation
relative standard deviation
Top-K (K=0 is sampled K1/K4)
Top-K, M32=0.2
Top-K, M32=0.5
Top-K, M32=0.8
Top-K, M32=0.9