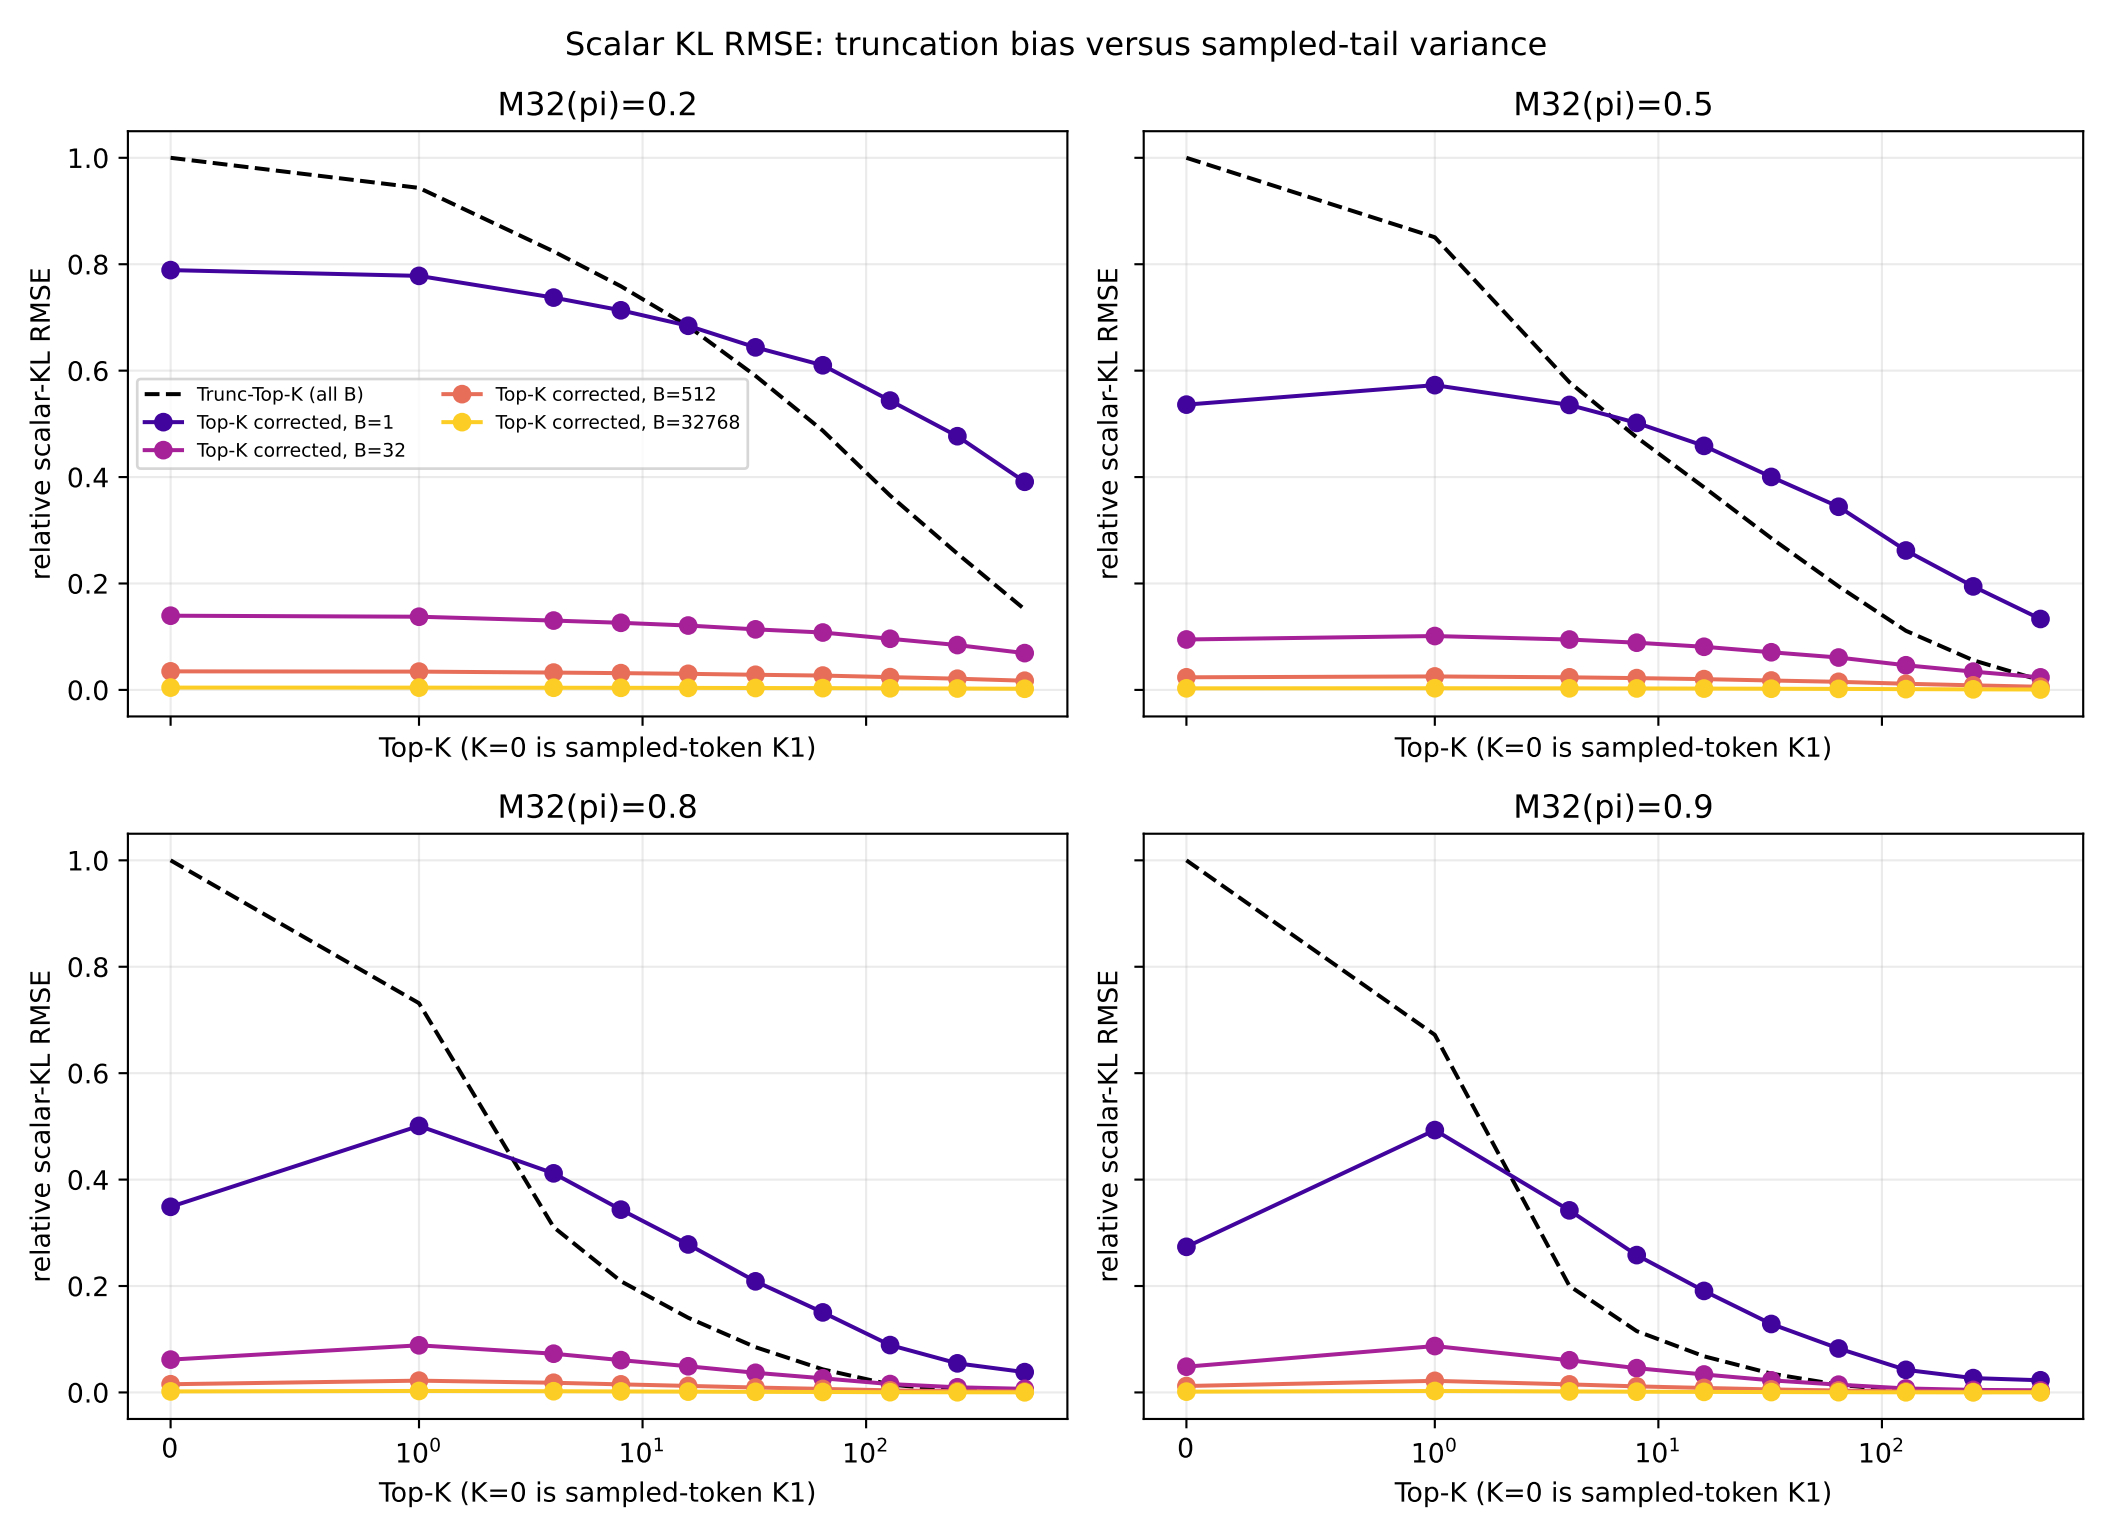
Scalar KL RMSE: truncation bias versus sampled-tail variance
M32(pi)=0.2
M32(pi)=0.5
M32(pi)=0.8
M32(pi)=0.9
relative scalar-KL RMSE
Top-K (K=0 is sampled-token K1)
Trunc-Top-K (all B)
Top-K corrected, B=1
Top-K corrected, B=32
Top-K corrected, B=512
Top-K corrected, B=32768
0.0
0.2
0.4
0.6
0.8
1.0
0
10^0
10^1
10^2