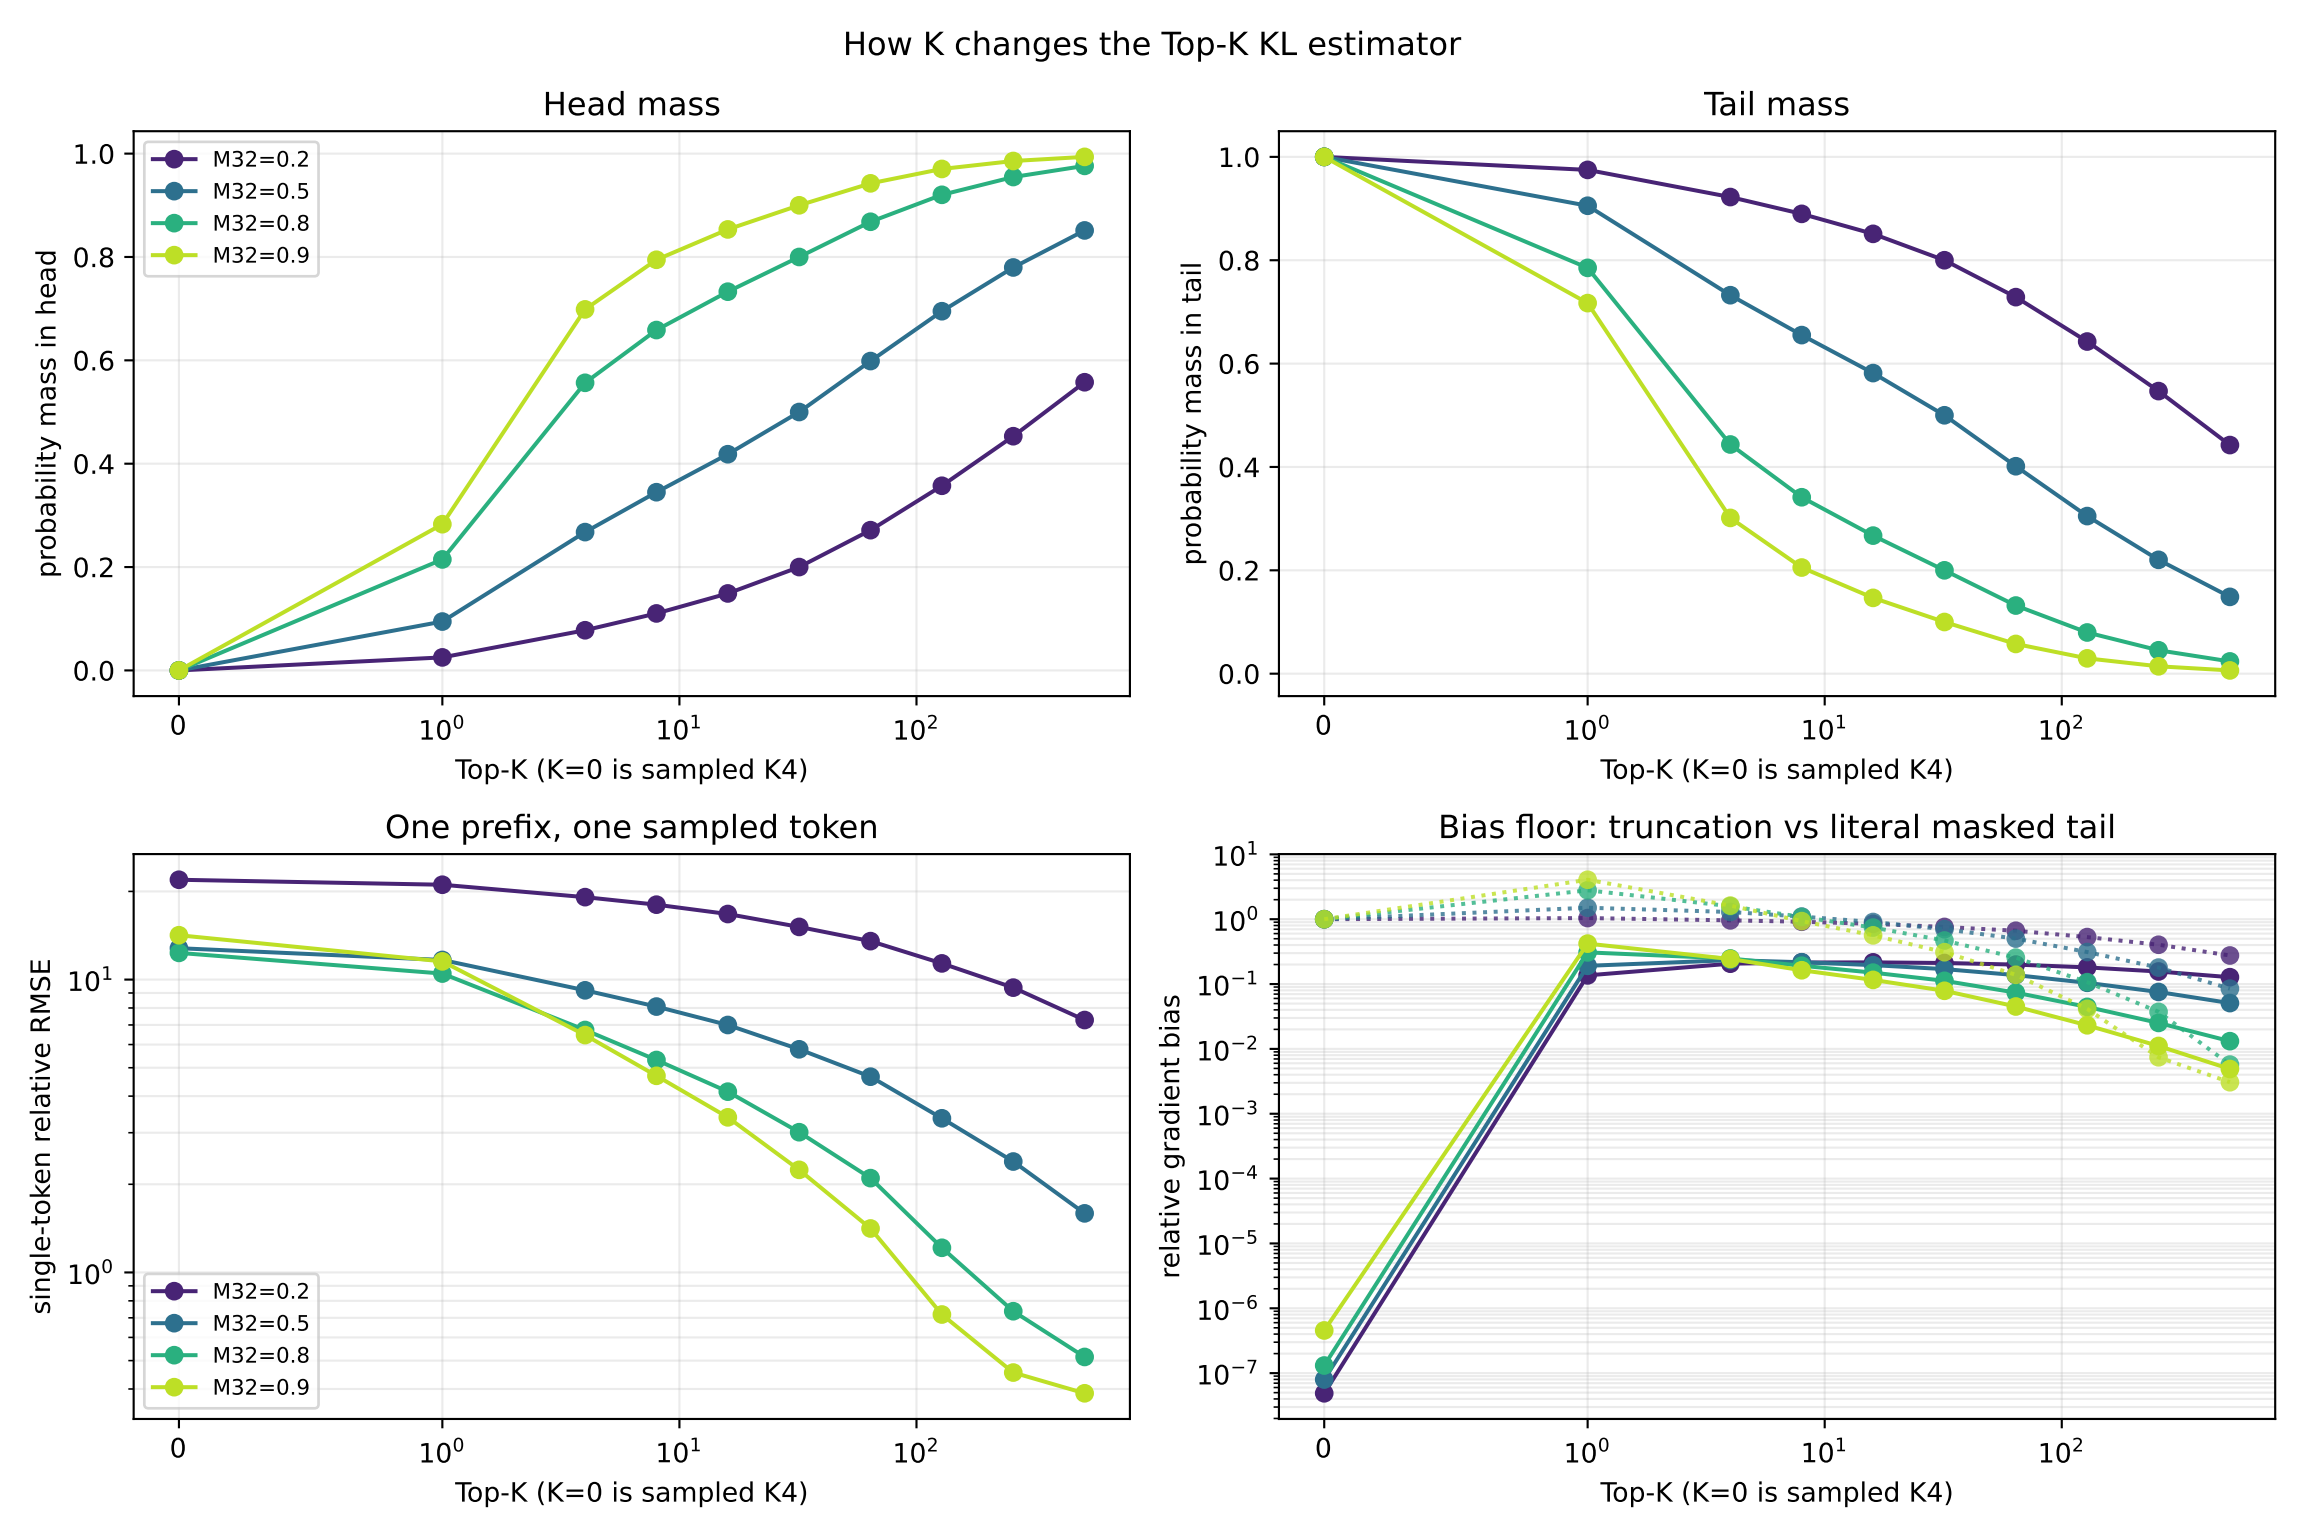
How K changes the Top-K KL estimator
Head mass
probability mass in head
Top-K (K=0 is sampled K4)
M32=0.2
M32=0.5
M32=0.8
M32=0.9
Tail mass
probability mass in tail
Top-K (K=0 is sampled K4)
One prefix, one sampled token
single-token relative RMSE
Top-K (K=0 is sampled K4)
Bias floor: truncation vs literal masked tail
relative gradient bias
Top-K (K=0 is sampled K4)