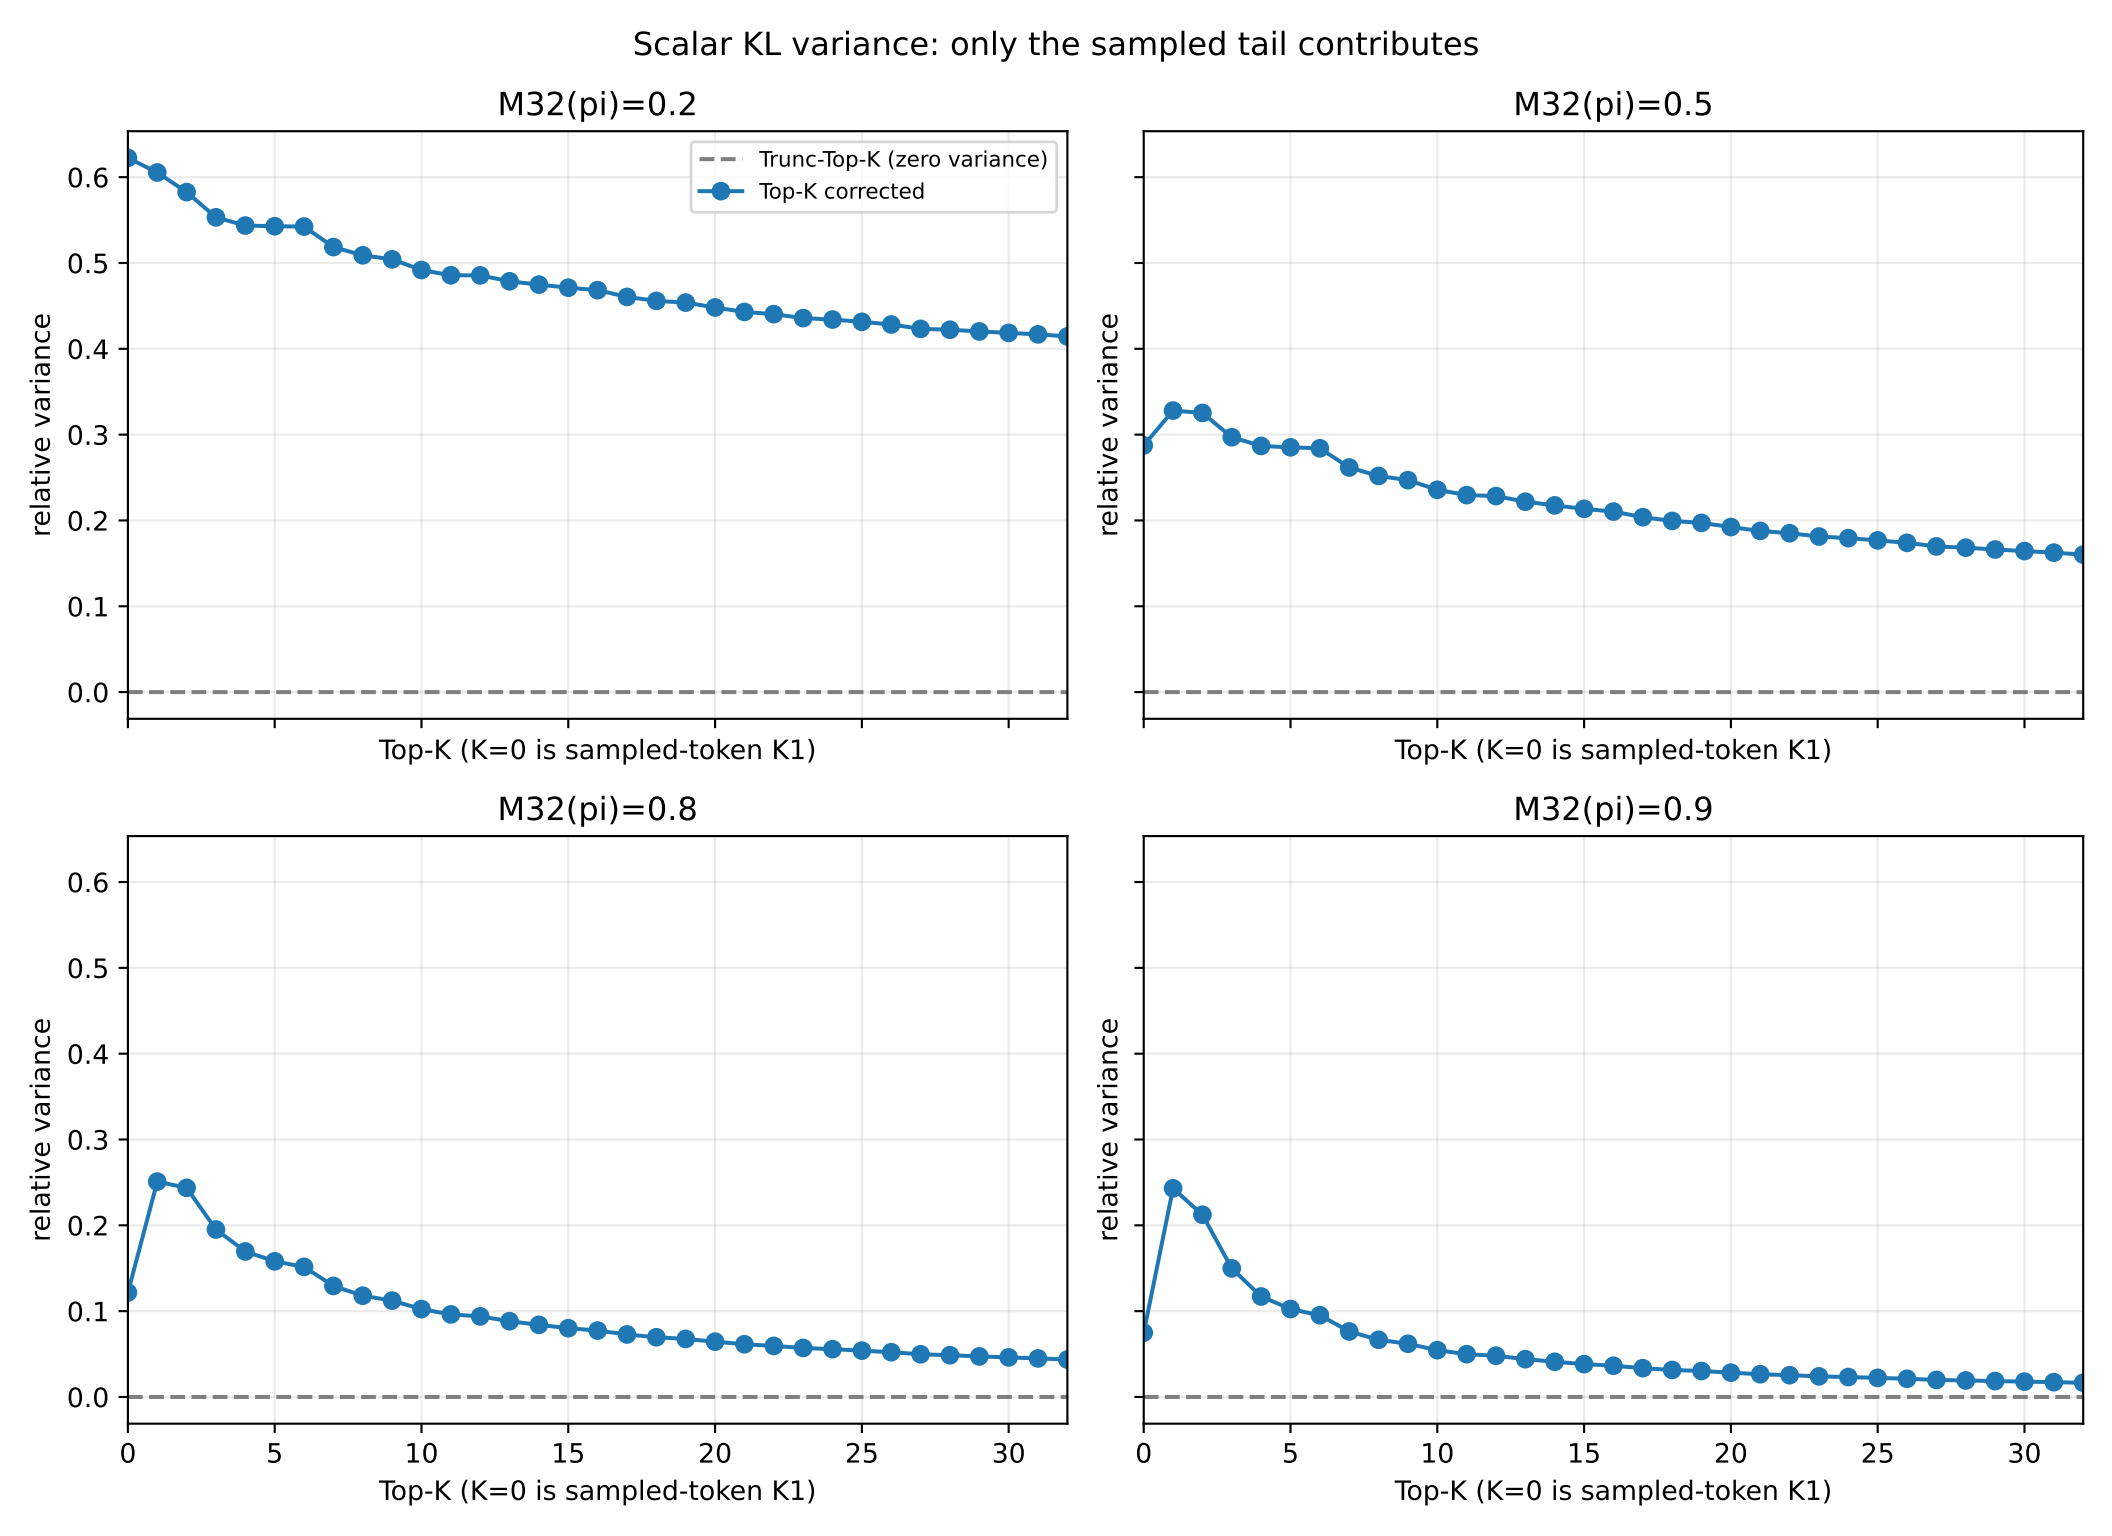
Scalar KL variance: only the sampled tail contributes
M32(pi)=0.2
M32(pi)=0.5
M32(pi)=0.8
M32(pi)=0.9
Trunc-Top-K (zero variance)
Top-K corrected
relative variance
Top-K (K=0 is sampled-token K1)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0
5
10
15
20
25
30