Seed 42, FP32 sampled BPO, frozen 64 x 256-token pool, minibatch 8
Independent 4096-position KL audits; dashed vertical line: descriptive step-60 anchor

Updates 0--600 | Updates 0--100 | Updates 500--600

Train KL(p || teacher)

Held-out KL(p || teacher)

Pre-clip grad norm
raw + trailing-20 median

Actual applied LR (zero retained)

Optimizer updates

