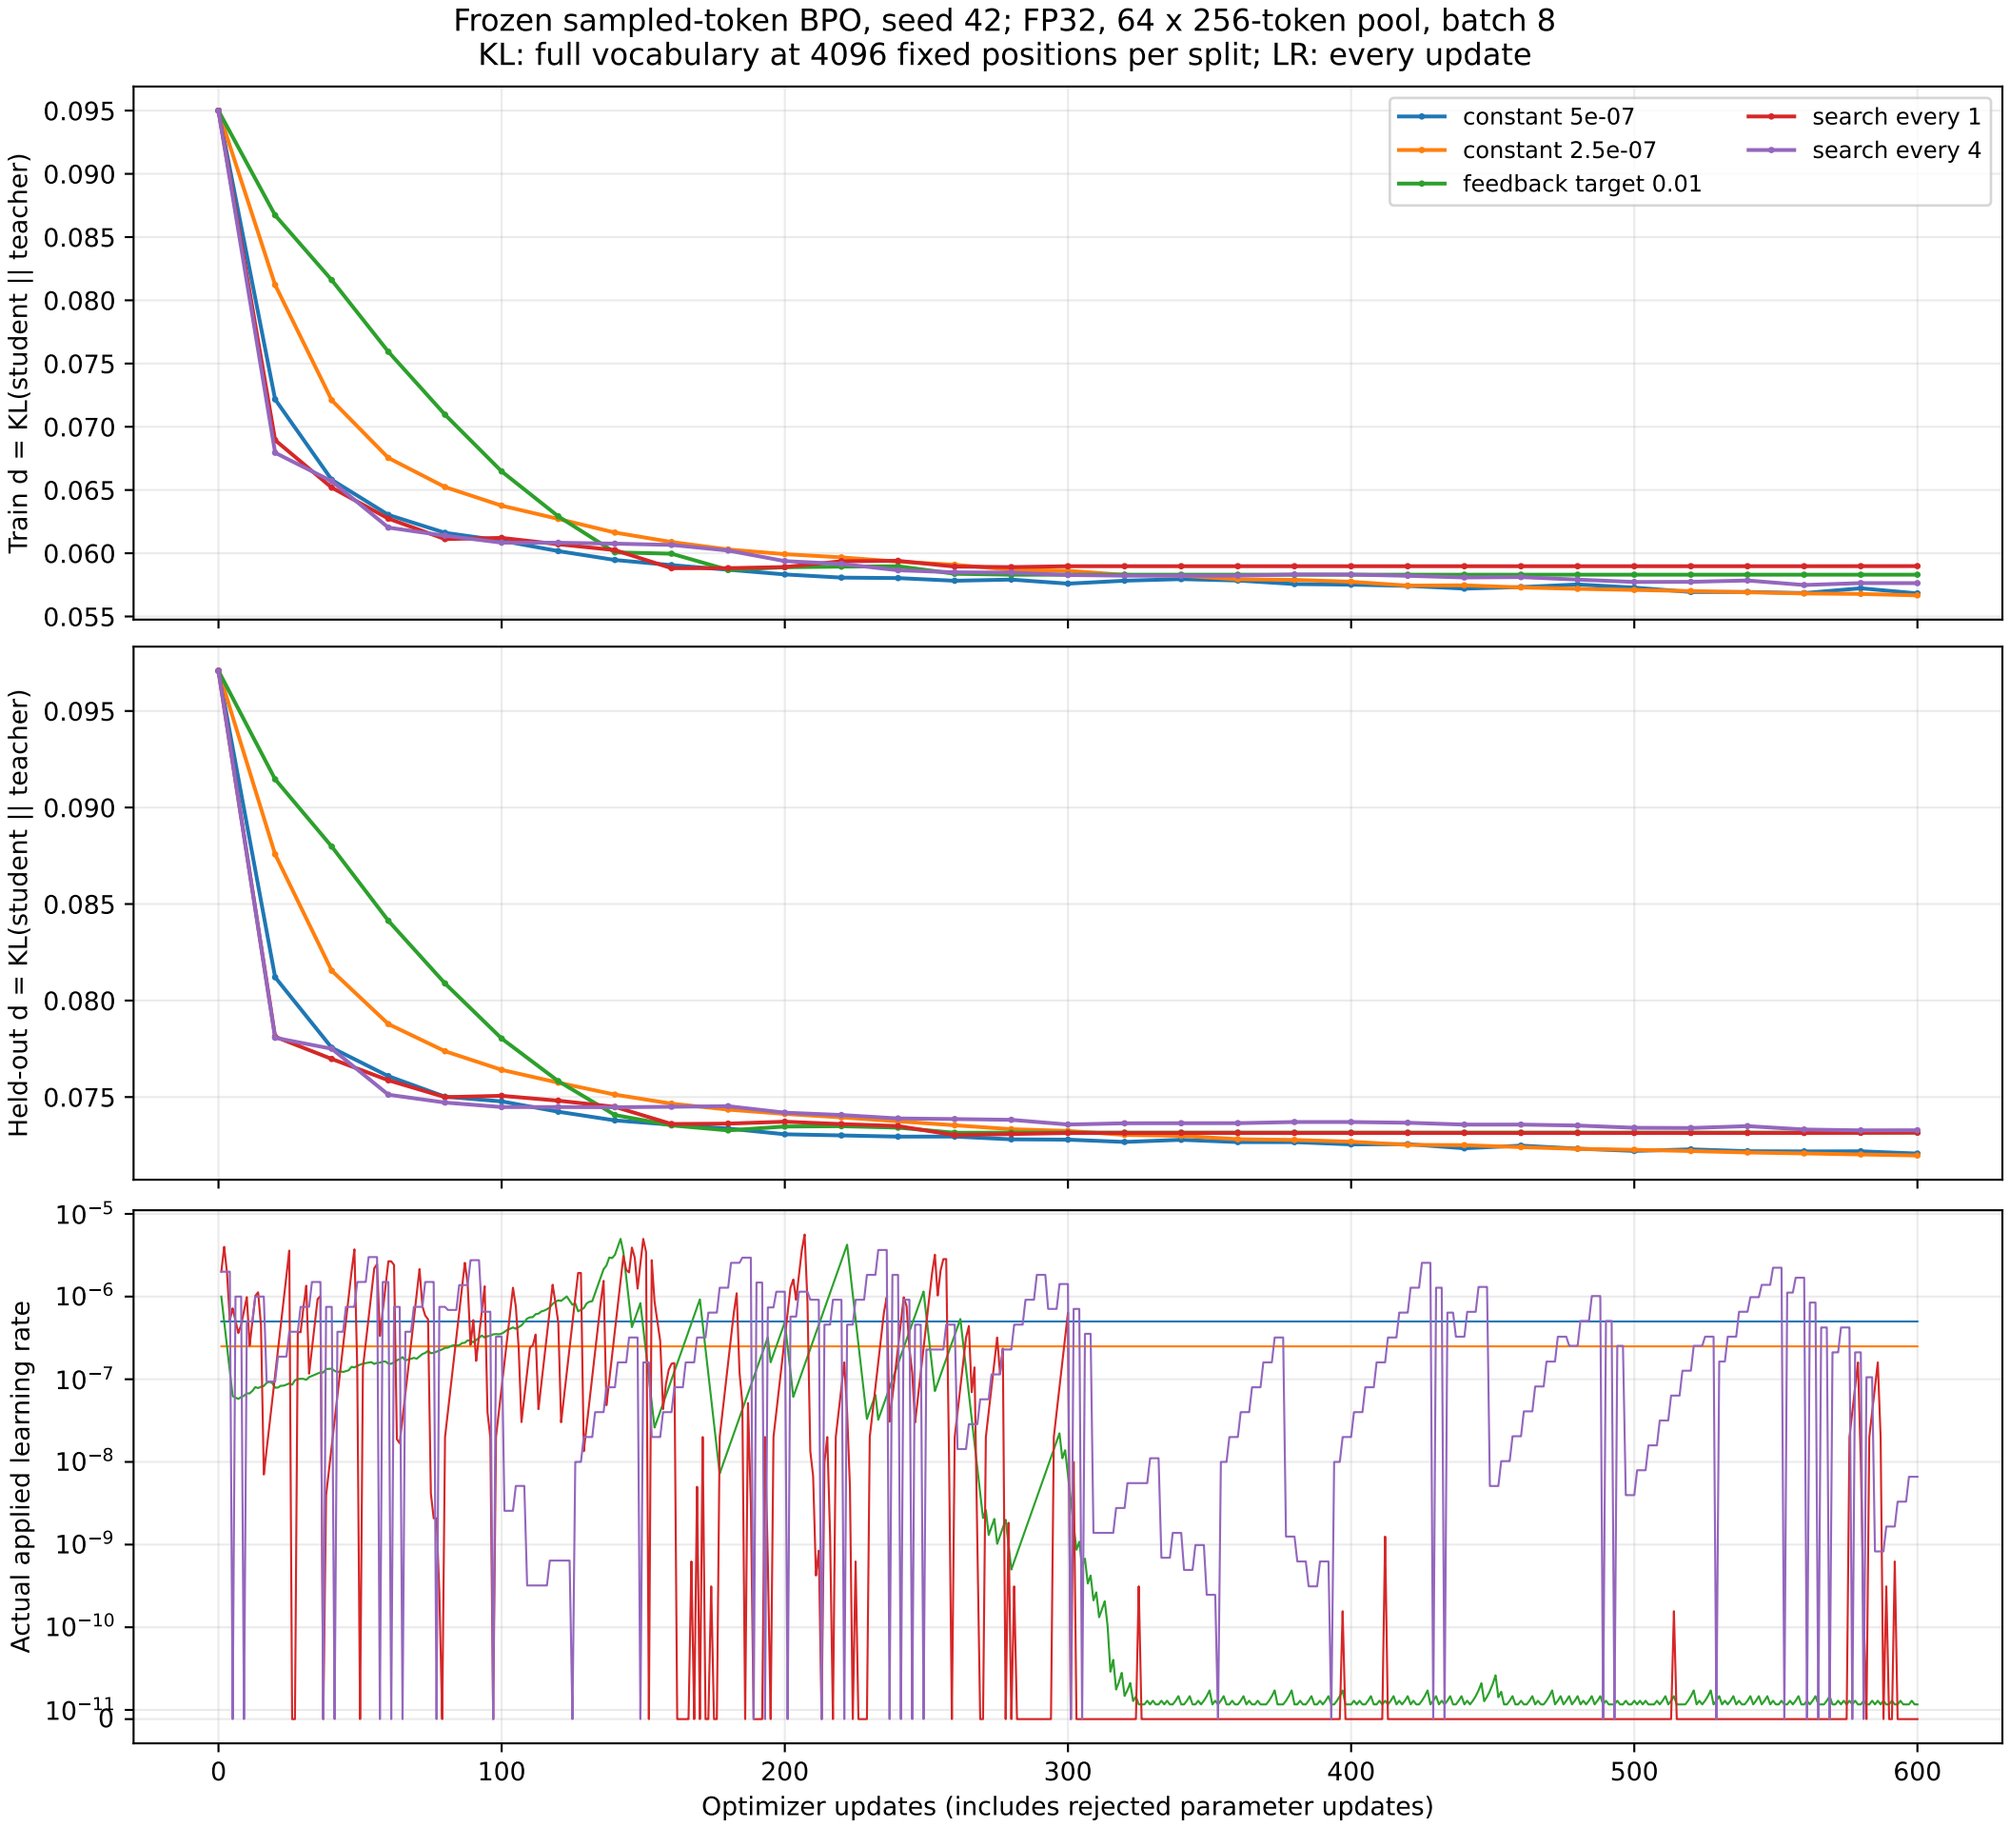
Frozen sampled-token BPO, seed 42; FP32, 64 x 256-token pool, batch 8
KL: full vocabulary at 4096 fixed positions per split; LR: every update
constant 5e-07
constant 2.5e-07
feedback target 0.01
search every 1
search every 4
Train d = KL(student || teacher)
Held-out d = KL(student || teacher)
Actual applied learning rate
Optimizer updates (includes rejected parameter updates)