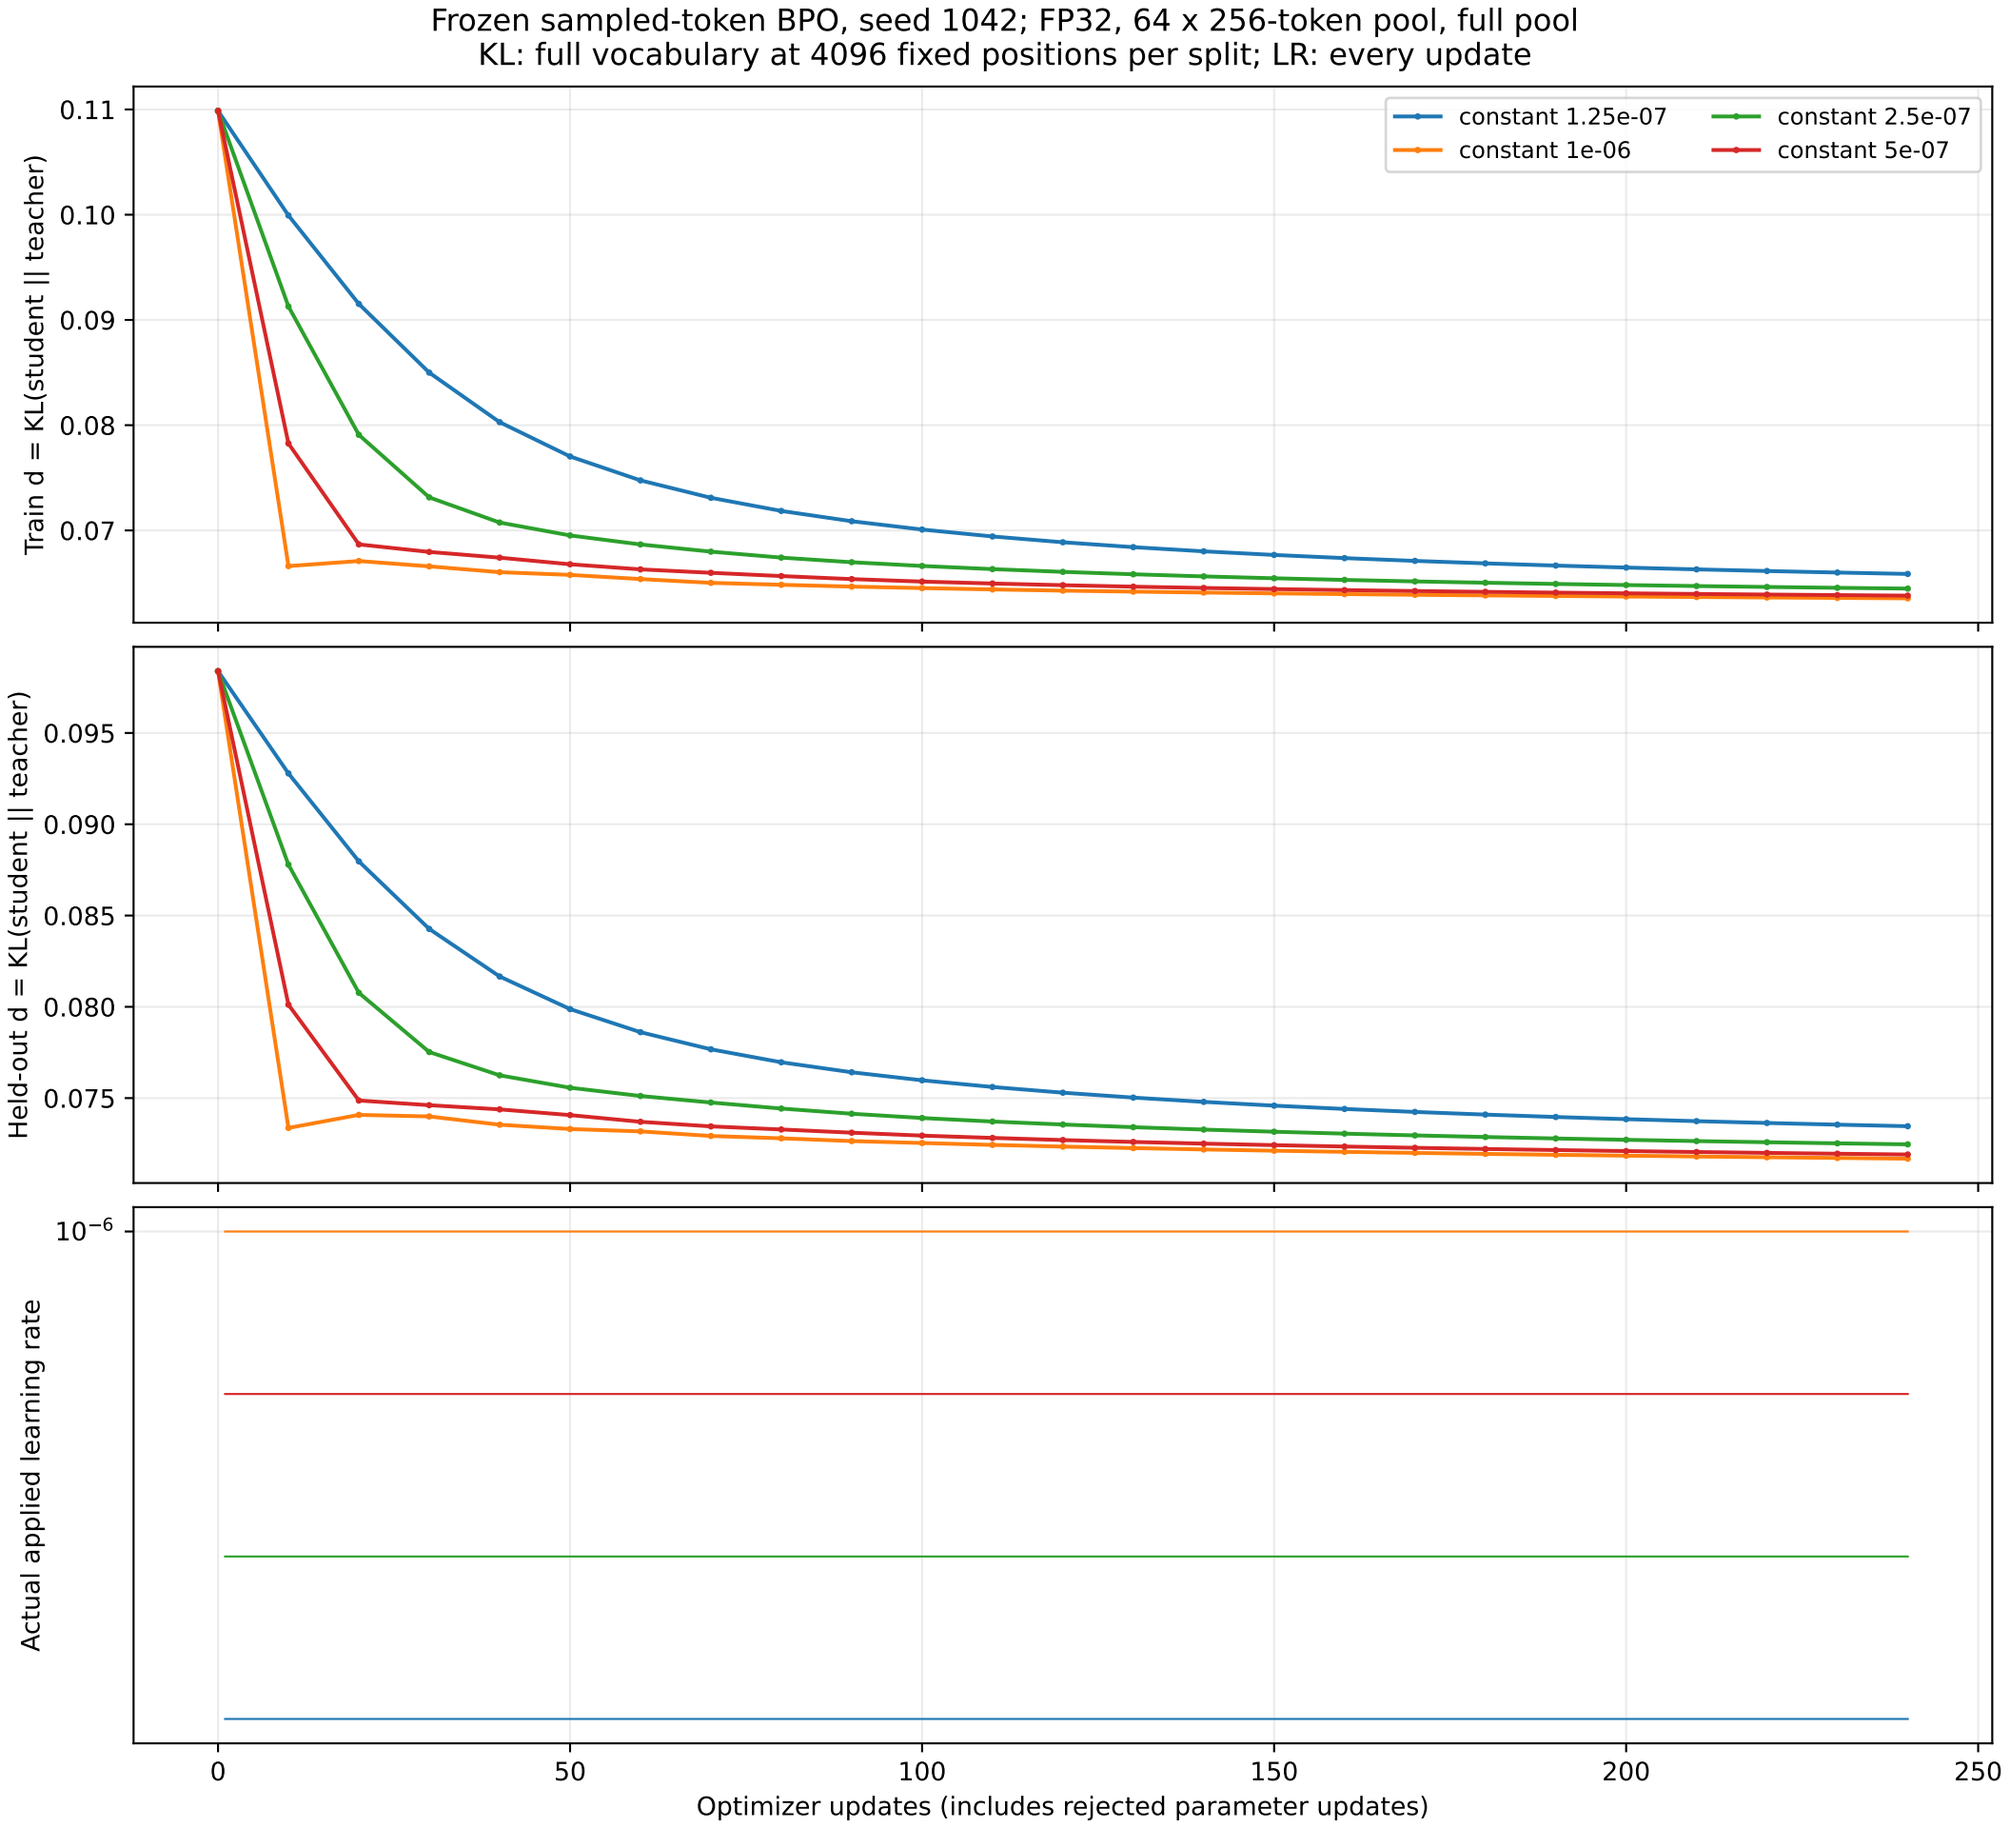
Frozen sampled-token BPO, seed 1042; FP32, 64 x 256-token pool, full pool
KL: full vocabulary at 4096 fixed positions per split; LR: every update
constant 1.25e-07
constant 2.5e-07
constant 1e-06
constant 5e-07
Train d = KL(student || teacher)
Held-out d = KL(student || teacher)
Actual applied learning rate
10^{-6}
Optimizer updates (includes rejected parameter updates)